

CZ4041/CE4041: Machine Learning

Lesson 11: Density Estimation

Kelly KE

School of Computer Science and Engineering,
NTU, Singapore


Acknowledgements: some figures are adapted from the lecture notes of Jason J. Corso (SUNY @ Buffalo). Slides are modified from the version prepared by Dr. Sinno Pan.

Density Estimation

- Density estimation aims to estimate an unobservable underlying probability density function based on observed data
- Denote by $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ the set of observed data points, drawn from an unknown $p(\mathbf{x})$,
$$\mathbf{x}_i \sim p(\mathbf{x}), \text{ for } i = 1, 2, \dots, N$$

The goal is to estimate the probability density function $p(\mathbf{x})$

Probability Density Function



$$P(\mathbf{x}^*) = \int_{x^*}^{x^* + \epsilon} p(x) dx \approx p(\mathbf{x}^*) \times \epsilon$$

Why Density Estimation

- Recall Naïve Bayes Classifiers

$$y_i = \arg \max_c P(\mathbf{x}|y = c)P(y = c)$$

$$= \arg \max_c \prod_{i=1}^d P(x_i|y = c) P(y = c)$$

- Other advanced ML topics
 - One-class classification
 - Covariate shift, transfer learning, etc.

Density Estimation Approaches

- Parametric density estimation
 - Assume a form for $p(\mathbf{x}; \boldsymbol{\theta})$, defined up to parameters, $\boldsymbol{\theta}$
 - E.g., Gaussian distribution $\mathcal{N}(\mu, \sigma^2)$, $\boldsymbol{\theta} = \{\mu, \sigma^2\}$
 - Estimate $\boldsymbol{\theta}$ from the observed data points
 - Maximum Likelihood Estimation
- Nonparametric density estimation

The General Principle

- The observed data points $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ are assumed to be a sample of N random variables independent and identically distributed (i.i.d.)
- Identically distributed: for any $\mathbf{x}_i \in \mathcal{D}$, it is sampled from the same probability distribution
- Independent: all the data points $\mathbf{x}_i \in \mathcal{D}$ are independent events

Parametric Density Estimation

- Assume that $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ are drawn from some known probability density family, $p(\mathbf{x}; \boldsymbol{\theta})$, defined up to parameters, $\boldsymbol{\theta}$
$$\mathbf{x}_i \sim p(\mathbf{x}; \boldsymbol{\theta})$$
- To find $\boldsymbol{\theta}$ that makes sampling \mathbf{x}_i from $p(\mathbf{x}; \boldsymbol{\theta})$ as likely as possible
- Approach: maximum likelihood estimation

Maximum Likelihood Estimation

- Likelihood of parameter θ given sample \mathcal{D} :

$$l(\theta|\mathcal{D}) \triangleq p(\mathcal{D}; \theta)$$

- As $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ are i.i.d., the above likelihood is the product of the likelihoods of the individual data points

$$l(\theta|\mathcal{D}) \triangleq p(\mathcal{D}; \theta) = \prod_{i=1}^N p(\mathbf{x}_i; \theta)$$

- In MLE, we aim to find θ that makes \mathcal{D} the most likely to be drawn from. Mathematically, we aim to search for $\hat{\theta}$ such that

$$\hat{\theta} = \arg \max_{\theta} l(\theta|\mathcal{D})$$

MLE (cont.)

- Typically, we maximize the log-likelihood:


$$\mathcal{L}(\boldsymbol{\theta}|\mathcal{D}) \triangleq \ln l(\boldsymbol{\theta}|\mathcal{D})$$

- Why?

- The $\ln(\cdot)$ function converts the product into a sum


$$\ln l(\boldsymbol{\theta}|\mathcal{D}) = \ln p(\mathcal{D}; \boldsymbol{\theta}) = \ln \left(\prod_{i=1}^N p(\mathbf{x}_i; \boldsymbol{\theta}) \right) = \sum_{i=1}^N \ln p(\mathbf{x}_i; \boldsymbol{\theta})$$

- The $\ln(\cdot)$ function is a strictly increasing function, one can maximize the likelihood without changing the value where it takes its maximum



$$\hat{\boldsymbol{\theta}} = \arg \max_{\boldsymbol{\theta}} l(\boldsymbol{\theta}|\mathcal{D}) \Leftrightarrow \hat{\boldsymbol{\theta}} = \arg \max_{\boldsymbol{\theta}} \ln l(\boldsymbol{\theta}|\mathcal{D})$$

MLE Illustration

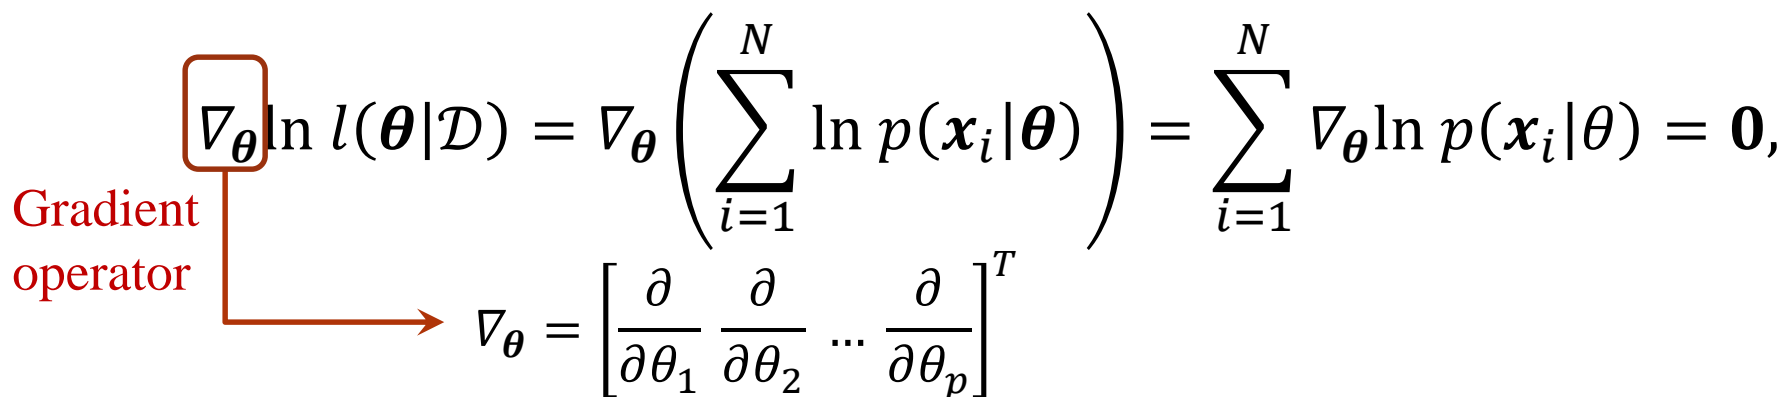


Solution of MLE

- Suppose $\boldsymbol{\theta}$ contains p parameters,

$$\boldsymbol{\theta} = [\theta_1 \ \theta_2 \ \dots \ \theta_p]^T$$

- As $\max_{\boldsymbol{\theta}} \ln l(\boldsymbol{\theta}|\mathcal{D})$ is an unconstrained optimization problem, to solve it, we first set to the derivative of $\ln l(\boldsymbol{\theta}|\mathcal{D})$ w.r.t. $\boldsymbol{\theta}$ to be zero

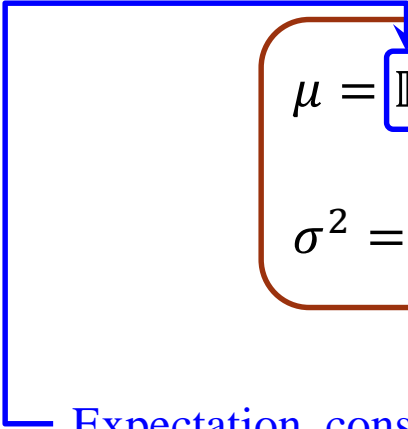

$$\nabla_{\boldsymbol{\theta}} \ln l(\boldsymbol{\theta}|\mathcal{D}) = \nabla_{\boldsymbol{\theta}} \left(\sum_{i=1}^N \ln p(\mathbf{x}_i|\boldsymbol{\theta}) \right) = \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} \ln p(\mathbf{x}_i|\boldsymbol{\theta}) = \mathbf{0},$$
$$\nabla_{\boldsymbol{\theta}} = \left[\frac{\partial}{\partial \theta_1} \ \frac{\partial}{\partial \theta_2} \ \dots \ \frac{\partial}{\partial \theta_p} \right]^T$$

- We then obtain a solution $\hat{\boldsymbol{\theta}}$ by solving the above system of p equations

Univariate Gaussian

- Suppose $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$. Each data point \mathbf{x}_i is of 1 dimension, and drawn from Gaussian distribution with unknown μ and σ^2 :

$$\mathbf{x}_i \sim p(\mathbf{x}; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\left(-\frac{(\mathbf{x}-\mu)^2}{2\sigma^2}\right)}$$


$$\mu = \mathbb{E}[\mathbf{x}] = \int p(\mathbf{x}; \mu, \sigma^2) \mathbf{x} d\mathbf{x}$$

$$\sigma^2 = \text{cov}(\mathbf{x}) = \mathbb{E}[(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{x} - \mathbb{E}[\mathbf{x}])^T]$$

Expectation, considering all possible values (infinite)

Univariate Gaussian (cont.)

- The log-likelihood is

$$\begin{aligned}\ln l(\boldsymbol{\theta}|\mathcal{D}) &= \sum_{i=1}^N \ln p(\mathbf{x}_i; \boldsymbol{\theta}) \\ p(\mathbf{x}; \mu, \sigma^2) &= \frac{1}{\sqrt{2\pi\sigma^2}} e^{\left(-\frac{(\mathbf{x}-\mu)^2}{2\sigma^2}\right)} \\ &= \sum_{i=1}^N \ln \left(\cancel{\exp} \left(-\frac{(\mathbf{x}_i - \mu)^2}{2\sigma^2} \right) \right) - \sum_{i=1}^N \ln(\sqrt{2\pi\sigma^2}) \\ &= -\frac{\sum_{i=1}^N (\mathbf{x}_i - \mu)^2}{2\sigma^2} - \frac{N}{2} \ln(2\pi\sigma^2) \\ &= -\frac{\sum_{i=1}^N (\mathbf{x}_i - \mu)^2}{2\sigma^2} - \frac{N}{2} \ln(2\pi) - \frac{N}{2} \ln \sigma^2\end{aligned}$$

Univariate Gaussian (cont.)

- The log-likelihood:

$$\ln l(\boldsymbol{\theta}|\mathcal{D}) = -\frac{\sum_{i=1}^N (\mathbf{x}_i - \mu)^2}{2\sigma^2} - \frac{N}{2} \ln(2\pi) - \frac{N}{2} \ln \sigma^2$$

- The derivative of the log-likelihood:

$$\begin{aligned} \nabla_{\boldsymbol{\theta}} \ln l(\boldsymbol{\theta}|\mathcal{D}) &= \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} \ln p(\mathbf{x}_i; \boldsymbol{\theta}) = \begin{bmatrix} \sum_{i=1}^N \nabla_{\mu} \ln p(\mathbf{x}_i; \boldsymbol{\theta}) \\ \sum_{i=1}^N \nabla_{\sigma^2} \ln p(\mathbf{x}_i; \boldsymbol{\theta}) \end{bmatrix} \\ &= \begin{bmatrix} \frac{1}{\sigma^2} \sum_{i=1}^N (\mathbf{x}_i - \mu) \\ \frac{\sum_{i=1}^N (\mathbf{x}_i - \mu)^2}{2(\sigma^2)^2} - \frac{N}{2\sigma^2} \end{bmatrix} \end{aligned}$$

Univariate Gaussian (cont.)

- By setting the derivative to be zero:

$$\begin{cases} \frac{1}{\sigma^2} \sum_{i=1}^N (x_i - \mu) = 0 \\ \frac{\sum_{i=1}^N (x_i - \mu)^2}{2(\sigma^2)^2} - \frac{N}{2\sigma^2} = 0 \end{cases}$$

- We have

$$\begin{cases} \hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i \\ \widehat{\sigma^2} = \frac{1}{N} \sum_{i=1}^N (x_i - \hat{\mu})^2 \end{cases}$$

Unbiased estimation

$$\mathbb{E}[\hat{\mu}] = \boxed{\mu} \leftarrow \text{True}$$


$$\mathbb{E}[\widehat{\sigma^2}] = \frac{N-1}{N} \boxed{\sigma^2}$$


Biased estimation

To correct bias

$$\begin{aligned} \widetilde{\sigma^2} &= \frac{N}{N-1} \widehat{\sigma^2} \\ &= \frac{1}{N-1} \sum_{i=1}^N (x_i - \hat{\mu})^2 \end{aligned}$$

Multivariate Gaussian

- Suppose $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$, and each data point \mathbf{x}_i is of d dimensions, and drawn from Gaussian distribution with unknown μ and Σ :


$$\mathbf{x}_i \sim p(\mathbf{x}; \mu, \Sigma) = \frac{1}{(2\pi)^{d/2}} \frac{1}{|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu)\right)$$


$|\Sigma|$: the determinant of Σ

$$\mu = \mathbb{E}[\mathbf{x}] = \int p(\mathbf{x}; \mu, \Sigma) \mathbf{x} d\mathbf{x}$$

$$\Sigma = \text{cov}(\mathbf{x}) = \mathbb{E}[(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{x} - \mathbb{E}[\mathbf{x}])^T]$$

Multivariate Gaussian (cont.)

- The log-likelihood of multivariate Gaussian is

$$\ln l(\boldsymbol{\theta}|\mathcal{D}) = \sum_{i=1}^N \ln p(\mathbf{x}_i; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = -\frac{Nd}{2} \ln(2\pi) - \frac{N}{2} \ln|\boldsymbol{\Sigma}| - \frac{1}{2} \sum_{i=1}^N (\mathbf{x}_i - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x}_i - \boldsymbol{\mu})$$

- The derivative of the log-likelihood w.r.t. $\boldsymbol{\mu}$ is

$$\frac{\partial}{\partial \boldsymbol{\mu}} \ln l(\boldsymbol{\theta}|\mathcal{D}) = \sum_{i=1}^N \boldsymbol{\Sigma}^{-1} (\mathbf{x}_i - \boldsymbol{\mu})$$

- By setting the derivative to zero, we have

$$\hat{\boldsymbol{\mu}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i$$

Unbiased estimation:

$$\mathbb{E}[\hat{\boldsymbol{\mu}}] = \boldsymbol{\mu}$$

Multivariate Gaussian (cont.)

- By computing the derivative of the log-likelihood w.r.t. Σ (each Σ_{ij} , i.e., the entry of the i -th row and the j -th column), and setting to be zero, we have

$$\hat{\Sigma} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \hat{\mu})(\mathbf{x}_i - \hat{\mu})^T$$

- The above estimation is biased!!!

$$\mathbb{E}[\hat{\Sigma}] = \frac{N-1}{N} \Sigma$$

- We can correct the bias by timing $\hat{\Sigma}$ with $\frac{N}{N-1}$:

$$\tilde{\Sigma} = \frac{1}{N-1} \sum_{i=1}^N (\mathbf{x}_i - \hat{\mu})(\mathbf{x}_i - \hat{\mu})^T \quad \text{Unbiased estimation}$$

Summary: MLE for Multivariate Gaussian

- Suppose $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$. Each data point \mathbf{x}_i is of d dimensions, and drawn from Gaussian distribution with unknown $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$:

$$\mathbf{x}_i \sim p(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2}} \frac{1}{|\boldsymbol{\Sigma}|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

- The unbiased MLE for $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ are

$$\left\{ \begin{array}{l} \hat{\boldsymbol{\mu}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i \\ \tilde{\boldsymbol{\Sigma}} = \frac{1}{N-1} \sum_{i=1}^N (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T \end{array} \right.$$

Note: in this module, when you are asked to estimate the mean vector and the covariate matrix of a sample, use the unbiased estimation





Density Estimation Approaches

- Parametric density estimation
- Nonparametric density estimation
 - Without assuming any forms for the underlying density
 - Assume that similar inputs have similar outputs: if \mathbf{x}_i and \mathbf{x}_j are similar, then $p(\mathbf{x}_i)$ and $p(\mathbf{x}_j)$ are similar
 - Approaches
 - Histogram Estimator
 - Naïve Estimator / Parzen Windows / Kernel Estimator
 - K -NN Estimator

Non-Parametric Density Estimation

- Assume that $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ are drawn from some unknown probability density $p(\mathbf{x})$
- To learn the estimator $\hat{p}(\mathbf{x})$ for $p(\mathbf{x})$
- We first focus on the univariate case, i.e., \mathbf{x}_i is 1-dimensional
- Note that the introduced approaches can be generalized to multivariate case easily

Histogram Estimator



Histogram Estimator (cont.)

- Simply partition \mathbf{x} into distinct bins of a fixed width Δ
- Count the number N_t of data points falling into bin t
- Turn this count into a normalized probability density via dividing by the total number of observed data points N and by the width Δ of the bins:

$$p_t = \frac{N_t}{N\Delta} \quad \text{Why divided by } \Delta?$$

- The model for the density $p(\mathbf{x})$ is constant over the width of each bin: find the bin where \mathbf{x} is in (e.g., bin t), then

$$\hat{p}(\mathbf{x}) = \frac{\#\{\mathbf{x}_i \mid \mathbf{x}_i \text{ in the same bin as } \mathbf{x}\}}{N\Delta} = p_t$$

Histogram Estimator (cont.)

- For a bin t , given a density function, the probability that a data instance falling into the bin t is

$$P_t = \int_{\Delta} p(\mathbf{x}) d\mathbf{x} = p_t(\mathbf{x})\Delta$$

- On the other hand,


$$P_t = \frac{\#\{\mathbf{x}_i \mid \mathbf{x}_i \text{ in bin } t\}}{N} = \frac{N_t}{N}$$

- Therefore:

$$p_t(\mathbf{x})\Delta = \frac{N_t}{N} \quad \longrightarrow \quad p_t(\mathbf{x}) = \frac{N_t}{N\Delta}$$

Histogram Estimator (cont.)

- Note: different bins may have different widths Δ_t in general, but in practice, we use the same width Δ




Histogram Estimator (cont.)

- Histogram density as a function of bin width Δ

When Δ is very small, the resulting density is quite spiky and hallucinates a lot of structure not present in the true density



The green curve is the underlying true density from which the data points were drawn



When Δ is very big, the resulting density is quite smooth and consequently fails to capture the bimodality of the true density



Analysis on Histogram Estimator

- Advantages:
 - Simple to evaluate and simple to use
 - One can throw away \mathcal{D} once the histogram is computed
 - Can be updated incrementally
- Disadvantages:
 - The estimated density has discontinuities due to the bin edges rather than any property of the underlying density
 - Scales poorly to multivariate cases: we would have m^d bins (hypercubes) if we divided each feature (dimension) in a d -dimensional space into m bins

Naïve Estimator: An Alternative

- In Histogram Estimator, besides Δ , we have to choose an origin x_0 as well, the bins are the intervals defined as


$$[x_0 + m\Delta, x_0 + (m + 1)\Delta)$$

 0, positive or negative integers

- The Naïve Estimator does not need to set an origin

$$\hat{p}(x) = \frac{\#\{x_i \mid x_i \text{ in the same bin as } x\}}{N\Delta}$$


Given x , use x as a center to create a bin with a length of Δ




$$\hat{p}(x) = \frac{\#\{x_i \mid x - \frac{\Delta}{2} \leq x_i < x + \frac{\Delta}{2}\}}{N\Delta}$$

Histogram v.s. Naïve Estimator

Histogram Estimator



Naïve Estimator



Naïve Estimator: An Alternative (cont.)

$$\hat{p}(\mathbf{x}) = \frac{\#\{\mathbf{x}_i \mid \mathbf{x} - \frac{\Delta}{2} \leq \mathbf{x}_i < \mathbf{x} + \frac{\Delta}{2}\}}{N\Delta}$$

The Naïve Estimator
can also be written as

$$\hat{p}(\mathbf{x}) = \frac{1}{N\Delta} \sum_{i=1}^N w\left(\frac{\mathbf{x} - \mathbf{x}_i}{\Delta}\right)$$




If \mathbf{x}_i is in the bin with
width Δ centered at \mathbf{x} ,
then the count is
increased by 1

$$u = \frac{\mathbf{x} - \mathbf{x}_i}{\Delta}$$

$$w(u) = \begin{cases} 1 & \text{if } -\frac{1}{2} \leq u < \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

Parzen windows

Naïve Estimator



Generalization to Multivariate


- Suppose the observed data points $\mathcal{D} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ are d -dimensional
- In a d -dimensional space, we define \mathcal{R} is a d -dimensional hypercube with h being the length of each edge. Then the volume of the hypercube is given by

$$V = h^d$$

- The windowing function can be defined as

$$u = \frac{\mathbf{x} - \mathbf{x}_i}{h} \quad w(\mathbf{u}) = \begin{cases} 1 & \text{if } -\frac{1}{2} \leq u_j < \frac{1}{2} \text{ for all } j \in \{1, 2, \dots, d\} \\ 0 & \text{otherwise} \end{cases}$$

E.g., $d = 2$



Defines a hypercube of length h centered at \mathbf{x} , if \mathbf{x}_i falls in the cube, then the count is increased by 1


Generalization to Multivariate (cont.)

- Hence, $w\left(\frac{\mathbf{x}-\mathbf{x}_i}{h}\right)$ is equal to unity if \mathbf{x}_i falls within the hypercube of volume V centered at \mathbf{x} , and is zero otherwise
- The density estimator can be written as

$$\hat{p}(\mathbf{x}) = \frac{\#\{\mathbf{x}_i \mid \mathbf{x}_i \text{ in the same hypercube as } \mathbf{x}\}}{NV}$$



$$\hat{p}(\mathbf{x}) = \frac{1}{NV} \sum_{i=1}^N w\left(\frac{\mathbf{x}-\mathbf{x}_i}{h}\right)$$



$$\hat{p}(\mathbf{x}) = \frac{3}{10h^2}$$

Parzen window
function can be a
kernel function –
Kernel estimator
(Appendix, optional)

K-Nearest Neighbor Estimator

- Recall

$$p(\mathbf{x}) = \frac{\overset{K_{\mathbf{x}}}{\#\{\mathbf{x}_i \mid \mathbf{x}_i \text{ in the same hypercube as } \mathbf{x}\}}}{VN}$$

- In the previous approaches, V (or Δ for univariate) is fixed for different queries \mathbf{x} 's
- The K -NN Estimator *adapts* the amount of smoothing to the *local* density of data, and the degree of smoothing is controlled by K , the number of neighbors

Consider K nearest neighbors of \mathbf{x}

$$p(\mathbf{x}) = \frac{K}{NV_{\mathbf{x}}}$$

The volume of the space centered at \mathbf{x} that exactly contains K nearest neighbors of \mathbf{x}


- Note that the number of nearest neighbors is typically much smaller than the number of training data points, i.e., $K \ll N$

Naïve Estimator v.s. K -NN Estimator

Univariate Case

Naïve Estimator


$$\hat{p}(x) = \frac{\#\{x_i \mid x - \frac{\Delta}{2} \leq x_i < x + \frac{\Delta}{2}\}}{N\Delta}$$



- Fix Δ , and check how many data points fall in the bin

K -NN Estimator

$$\hat{p}(x) = \frac{K}{N(2d_K(x))}$$



- Fix K , the number of observed data points to fall in the bin, and compute the bin size

***K*-Nearest Neighbor Estimator (cont.)**

- With a predefined K , for a particular data point \mathbf{x} ,
 1. Compute distance between \mathbf{x} and all the observed data, e.g., Euclidean distance $\|\mathbf{x} - \mathbf{x}_i\|_2$
 2. Sort the observed data points based on the distances in ascending order:

$$d_1(\mathbf{x}) \leq d_2(\mathbf{x}) \leq \cdots d_j(\mathbf{x}) \leq \cdots \leq d_N(\mathbf{x})$$

$d_1(\mathbf{x})$ is the distance of \mathbf{x} to the nearest observed instance

$d_j(\mathbf{x})$ is the distance of \mathbf{x} to the j -th nearest observed instance

3. The K -NN density estimate is

$$\hat{p}(\mathbf{x}) = \frac{K}{N V_{d_K(\mathbf{x})}}$$

A ball in d -dimensional Euclidean space




The volume of the d -ball of the radius $d_K(\mathbf{x})$ centered at \mathbf{x} . And $d_K(\mathbf{x})$ is the distance of \mathbf{x} to the K -th nearest observed instance

Look-up table for Volume of an n -ball

https://en.wikipedia.org/wiki/Volume_of_an_n-ball

Dimension	Volume of a ball of radius R	Radius of a ball of volume V
0	1	(all 0-balls have volume 1)
1	$2R$	$\frac{V}{2} = 0.5 \times V$
2	$\pi R^2 \approx 3.142 \times R^2$	$\frac{V^{\frac{1}{2}}}{\sqrt{\pi}} \approx 0.564 \times V^{\frac{1}{2}}$
3	$\frac{4\pi}{3} R^3 \approx 4.189 \times R^3$	$\left(\frac{3V}{4\pi}\right)^{\frac{1}{3}} \approx 0.620 \times V^{\frac{1}{3}}$
4	$\frac{\pi^2}{2} R^4 \approx 4.935 \times R^4$	$\frac{(2V)^{\frac{1}{4}}}{\sqrt{\pi}} \approx 0.671 \times V^{\frac{1}{4}}$
5	$\frac{8\pi^2}{15} R^5 \approx 5.264 \times R^5$	$\left(\frac{15V}{8\pi^2}\right)^{\frac{1}{5}} \approx 0.717 \times V^{\frac{1}{5}}$
6	$\frac{\pi^3}{6} R^6 \approx 5.168 \times R^6$	$\frac{(6V)^{\frac{1}{6}}}{\sqrt{\pi}} \approx 0.761 \times V^{\frac{1}{6}}$
7	$\frac{16\pi^3}{105} R^7 \approx 4.725 \times R^7$	$\left(\frac{105V}{16\pi^3}\right)^{\frac{1}{7}} \approx 0.801 \times V^{\frac{1}{7}}$
8	$\frac{\pi^4}{24} R^8 \approx 4.059 \times R^8$	$\frac{(24V)^{\frac{1}{8}}}{\sqrt{\pi}} \approx 0.839 \times V^{\frac{1}{8}}$
9	$\frac{32\pi^4}{945} R^9 \approx 3.299 \times R^9$	$\left(\frac{945V}{32\pi^4}\right)^{\frac{1}{9}} \approx 0.876 \times V^{\frac{1}{9}}$
10	$\frac{\pi^5}{120} R^{10} \approx 2.550 \times R^{10}$	$\frac{(120V)^{\frac{1}{10}}}{\sqrt{\pi}} \approx 0.911 \times V^{\frac{1}{10}}$
11	$\frac{64\pi^5}{10395} R^{11} \approx 1.884 \times R^{11}$	$\left(\frac{10395V}{64\pi^5}\right)^{\frac{1}{11}} \approx 0.944 \times V^{\frac{1}{11}}$
12	$\frac{\pi^6}{720} R^{12} \approx 1.335 \times R^{12}$	$\frac{(720V)^{\frac{1}{12}}}{\sqrt{\pi}} \approx 0.976 \times V^{\frac{1}{12}}$
13	$\frac{128\pi^6}{135135} R^{13} \approx 0.911 \times R^{13}$	$\left(\frac{135135V}{128\pi^6}\right)^{\frac{1}{13}} \approx 1.007 \times V^{\frac{1}{13}}$
14	$\frac{\pi^7}{5040} R^{14} \approx 0.599 \times R^{14}$	$\frac{(5040V)^{\frac{1}{14}}}{\sqrt{\pi}} \approx 1.037 \times V^{\frac{1}{14}}$
15	$\frac{256\pi^7}{2027025} R^{15} \approx 0.381 \times R^{15}$	$\left(\frac{2027025V}{256\pi^7}\right)^{\frac{1}{15}} \approx 1.066 \times V^{\frac{1}{15}}$
n	$V_n(R)$	$R_n(V)$

K -NN Estimator



Thank you!

Appendix: Kernel Estimator

- To get a smooth estimate, we use a smooth weight function, *kernel function*, e.g., the Gaussian kernel
- For the univariate case, i.e., each data point is 1-dimensional, the Gaussian kernel is defined as

$$k(u) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{u^2}{2}\right)$$

- The Kernel Estimator is computed via

$$\hat{p}(x) = \frac{1}{N\Delta} \sum_{i=1}^N k\left(\frac{x - x_i}{\Delta}\right)$$

Kernel Estimator (cont.)

- For the multivariate case, i.e., each data point is d -dimensional, the Gaussian kernel is defined as

$$k(\mathbf{u}) = \frac{1}{(2\pi)^{d/2}} \exp\left(-\frac{\|\mathbf{u}\|_2^2}{2}\right)$$

- The Kernel Estimator is computed via

$$\hat{p}(\mathbf{x}) = \frac{1}{NV} \sum_{i=1}^N k\left(\frac{\mathbf{x} - \mathbf{x}_i}{h}\right)$$

Kernel Estimator

Optional

